

Server and cache mode expose the real Flashcache gain

Prompt-eval reduction across seven gpt-oss-20b Printy turns; larger is better.

- Per-request cold
- Persistent cold
- Hot cache
- Session tail
- Boundary telemetry
- Version cached
- Session full prompt

