

Warm runs are much cheaper than restarted runs

Ollama gpt-oss:20b on the Printy changed-tail workflow

