

Hot-cache savings grow with reusable context

gpt-oss-20b GGUF, persistent server, prewarmed stable-prefix slots

